Benchmarking anomaly detection techniques in the fraud detection domain*

*Note: Sub-titles are not captured in Xplore and should not be used

1st Leonardo Brito
Engineering and Data Science
Instituto Superior Técnico
Lisbon, Portugal
leonardo.amado.brito@tecnico.ulisboa.pt

Company Supervisor: Jacopo Bono
Data Scientist
Feedzai
Porto, Portugal
jacopo.bono@feedzai.com

Abstract—This document is a model and instructions for LaTeX. This and the IEEEtran.cls file define the components of your paper [title, text, heads, etc.]. *CRITICAL: Do Not Use Symbols, Special Characters, Footnotes, or Math in Paper Title or Abstract.

Index Terms—component, formatting, style, styling, insert

I. Introduction

In the realm of digital transactions, the significance of detecting fraud at its earliest is paramount. Supervised learning models are often the preferred choice, renowned for their accuracy, but they hinge on the availability of labeled data. This dependency can introduce operational delays, leaving a window where transactions might be exposed to unchecked fraudulent activities.

To address this gap, there's a growing interest in leveraging unsupervised learning. Without the need for labeled data, unsupervised anomaly detection techniques offer a promising avenue, especially when one considers that fraud typically surfaces as statistical anomalies amidst legitimate transactions.

Drawing from open-source datasets, this report sets out to critically evaluate a range of unsupervised anomaly detection methods in the context of fraud detection. Our approach is methodical: starting with a detailed data exploration, we then benchmark against a supervised baseline, and finally, analyze the performance of both classical and deep learning unsupervised models. The overarching aim is to ascertain the potential of unsupervised techniques in offering interim protection against fraud during periods when supervised models are not yet feasible.

II. Exploratory Data Analysis

A. European cardholders's transactions

This dataset represents credit card transactions made by European cardholders in September 2013 over a span of two days.

Out of the 284,807 transactions recorded, 492 were fraudulent, making up a mere 0.172% of the total transactions. Given the disparity in numbers, the dataset is notably imbalanced.

The dataset consists mainly of numerical variables derived from a Principal Component Analysis (PCA) transformation. To maintain confidentiality, the original features and further details about the data are not disclosed. The variables V1, V2, through V28 are the principal components derived from the PCA. The only exceptions that haven't undergone PCA are 'Time' and 'Amount'. The 'Time' feature indicates the seconds that have passed between each transaction and the first one in the dataset. On the other hand, 'Amount' denotes the transaction value, which can be useful for context-sensitive learning approaches. The 'Class' feature is the target variable, where a value of 1 indicates a fraudulent transaction, and 0 signifies a legitimate one.

There is no null or missing values of the dataset, but was 1854 observations duplicated, which were removed because they don't add any value to the analysis.

TABLE I
SUMMARY STATISTICS FOR FRAUD AND VALID TRANSACTIONS

Statistic	Fraud	Valid
Count	492	284315
Mean	122.21	88.29
Standard Deviation	256.68	250.11
Minimum	0.0	0.0
25% Quantile	1.0	5.65
50% Quantile (Median)	9.25	22.0
75% Quantile	105.89	77.05
Maximum	2125.87	25691.16

B. Analysis of Transaction Amounts for Fraudulent and Valid Transactions

The table presents a summary of transaction amounts for both fraudulent and valid categories. A notable observation is the disparity in the count of transactions: while there are only 492 fraudulent transactions, there are a staggering 284,315

valid ones. This vast difference highlights the class imbalance inherent in the dataset.

When examining the transaction amounts, fraudulent transactions have a mean value of approximately $122.21, which is slightly higher than the valid transactions' mean of $88.29. Despite this, the maximum fraudulent transaction ($2125.87) is significantly lower than the highest valid transaction, which reaches a substantial $25,691.16.

The spread of transaction amounts, as represented by the standard deviation, is roughly similar for both classes, hovering around the $250 mark. The median value for valid transactions ($22.00) is more than double that of fraudulent transactions ($9.25), indicating that half of the fraudulent transactions are below this amount.

While fraudulent transactions are comparatively rarer, their amounts can vary significantly, sometimes even surpassing the typical amounts seen in valid transactions. It underscores the importance of leveraging advanced techniques to detect such anomalies amidst the vast sea of valid transactions.

1) Imbalanced Dataset: An imbalanced (or unbalanced) dataset refers to a situation where the number of observations belonging to one class is significantly lower than those belonging to the other classes. In the context of a binary classification, which consists of two classes, an imbalanced dataset typically has a disproportionate ratio of observations in one class compared to the other. This problem can lead to:

- **Performance Deception:** Traditional metrics like accuracy can be misleading. A naive classifier predicting only the majority class would still achieve a very high accuracy due to the imbalance.
- **Model Bias:** Many machine learning models might exhibit a bias towards the majority class, often ignoring the minority class.
- **Decreased Predictive Power:** The minority class, in this case, the class of higher interest might not be predicted well, leading to a higher number of false negatives.

Given that our dataset has a class with 99.83% representation, it's highly imbalanced. I will present later on this report some considerations and how to deal with this problem.

subsubsectionCorrelation Matrix The fact the dataset is unbalanced can lead to misleading results when using correlation metrics:

- In highly imbalanced datasets, even small patterns in the minority class can result in seemingly strong correlations. This can lead to the model finding relationships that aren't generalizable.
- Correlation measures in an imbalanced dataset might be dominated by the majority class, potentially overshadowing meaningful relationships present in the minority class.
- Also, because the minority class has fewer data points, there's a higher risk of not detecting a relationship (false negative) when one might exist.

Due to the problem I presented, Traditional correlation metrics may not be suitable for imbalanced datasets. In this binary classification problem with a highly imbalanced target,

a high Pearson correlation coefficient with a predictor might be misleading.

2) Distributions: Due to the PCA transformation, the features are not directly interpretable. However, we can still examine the distributions of the features to gain some insights. The distributions of the features for both fraudulent and valid transactions are shown in Figure **??**. The goal was to understand the distribution of the features for both classes. Several features exhibit distinct distribution patterns when comparing the two Class categories: V4 and V11 show clear segregation between Class 0 and Class 1, suggesting high discriminatory power. Features V12, V14, and V18 display partial separation, indicating moderate distinction between classes. V1, V2, V3, and V10 each present a somewhat unique distribution, whereas V25, V26, and V28 appear to have overlapping distributions for both classes. Generally speaking, except for the 'Time' and 'Amount' features, the distribution of features for genuine transactions (Class = 0) tends to be centralized around 0, often with an extended tail at one end. Conversely, distributions for fraudulent transactions (Class = 1) are typically asymmetrical, reflecting a deviation from the norm.


Fig. 1. Example of a figure caption.

C. Variance

The majority of the data is the result of a PCA transformation - in this case was to **guarantee anonymity of the clients** - several unique characteristics and considerations come into play:

1) **Orthogonality:** The principal components resulting from PCA are orthogonal (uncorrelated). Thus, the correlation matrix of these components should be a diagonal matrix with ones on the diagonal (or very close to this in practice due to numerical precision).

2) **Variance Explained:** One of the key aspects of PCA is the amount of variance explained by each principal component. The first few components typically capture the majority of the variance in the dataset, while the latter components capture less and less variance.

Given these considerations, the variation analysis that makes the most sense for PCA-transformed data includes:

1) **Variance Explained:**

- **Scree Plot:** A plot showing the fraction of total variance explained by each principal component. This helps in determining how many components to retain for further analysis.

PCA is a dimensionality reduction technique that seeks to identify axes in data that maximize variance. The method involves computing the eigenvalues and eigenvectors of the dataset's covariance matrix. The eigenvectors, termed principal components, determine the direction of the new feature space, while the eigenvalues define their magnitude, i.e., the variance in those directions.

1) **Eigenvalues and Variance**: The eigenvalue associated with each principal component signifies the variance along that component. In PCA, these components are ordered by descending eigenvalues. This means the first principal component encapsulates the largest variance in the dataset.
2) **Orthogonality**: Every principal component is orthogonal to every other, implying they are uncorrelated. Hence, each subsequent component captures the direction of maximum variance not represented by the preceding components.
3) **Maximizing Variance**: The first principal component (often denoted as $PC1$) represents the direction in the original feature space capturing the utmost variance. The second principal component (often $PC2$) is orthogonal to $PC1$ and represents the second-highest variance, and so forth.

Consequently, when PCA-transformed features are acquired (like 'V1', 'V2', 'V3', etc.), they are inherently ordered by the variance they represent from the original dataset, with 'V1' representing the most and the final component the least.

III. Data Preprocessing

A. Linear VS Non-Linear Scaling

In summary, linear scaling methods are generally the first choice due to their simplicity and effectiveness in many scenarios. However, if you're dealing with outliers, skewed data, or specific nonlinear relationships, you might benefit from non-linear scaling methods. It's crucial to understand both your data and model requirements when making this decision.

B. Metrics used

F1 Score Explanation: The F1 score is a measure of a model's accuracy in a binary classification task, which considers both precision and recall of the test to compute the score. The F1 score can be understood as a weighted average (harmonic mean) of the precision and recall, where an F1 score reaches its best value at 1 (perfect precision and recall) and worst at 0.

Key Concepts

Precision: This is the number of true positive results divided by the number of all positive results, including those not identified correctly. It reflects the model's ability to identify only the relevant objects.

$$Precision = \frac{TP}{TP + FP}$$

Recall (Sensitivity or True Positive Rate): This is the number of true positive results divided by the number of all samples that should have been identified as positive. It reflects the ability of the model to find all the relevant cases within a dataset.

$$Recall = \frac{TP}{TP + FN}$$

F1 Score: This is the harmonic mean of precision and recall. The harmonic mean is used rather than the arithmetic mean as it punishes extreme values more. A model with perfect precision and recall will have an F1 score of 1. A model with a perfect precision but poor recall, or vice versa, will have an F1 score closer to 0.

$$F1 = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall}$$

Alternatively:

$$F1 = \frac{2TP}{2TP + FP + FN}$$

When to Use F1 Score

- Use the F1 Score when you seek a balance between precision and recall, especially if there is an uneven class distribution (large number of actual negatives).
- The F1 Score is utilized when both false positives and false negatives are equally costly, and it provides a single metric that summarizes model performance in a way that values precision and recall equally.

Considerations

- While the F1 Score is better than accuracy for imbalanced classes, in extremely imbalanced situations, even F1 might not be a perfect measure.
- In scenarios where one type of error has a significantly different consequence than the other, you might need to focus more on improving either recall or precision, and thus a different metric might be more appropriate.

REFERENCES

REFERENCES

[1] G. Eason, B. Noble, and I. N. Sneddon, "On certain integrals of Lipschitz-Hankel type involving products of Bessel functions," Phil. Trans. Roy. Soc. London, vol. A247, pp. 529–551, April 1955.
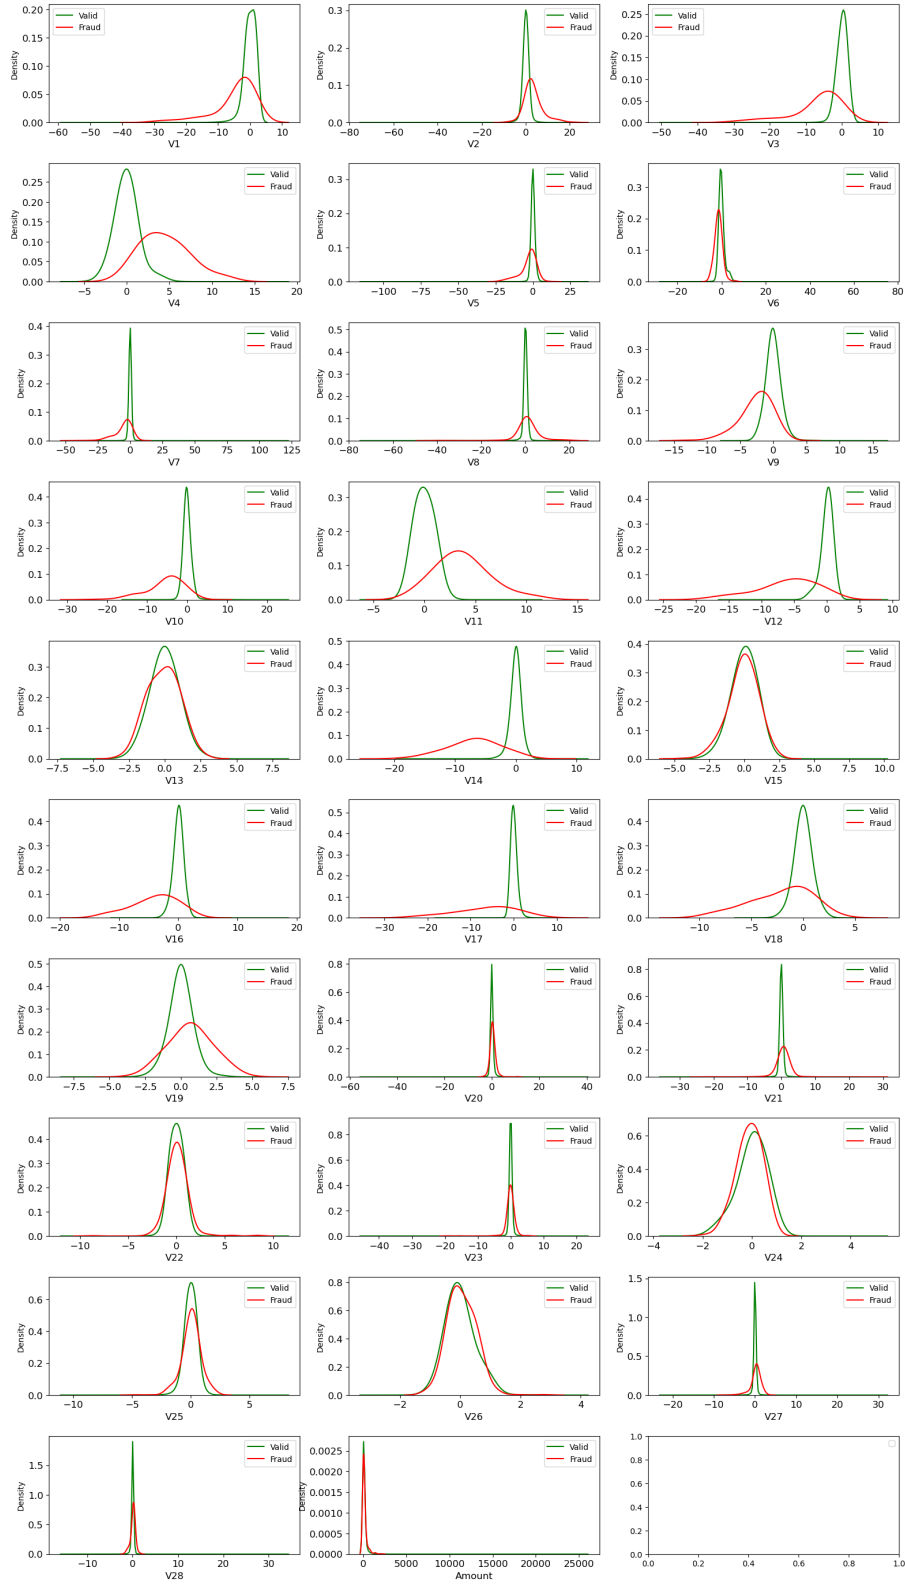
APPENDIX

Fig. 2. Distributions of all features for fraud and valid transactions.